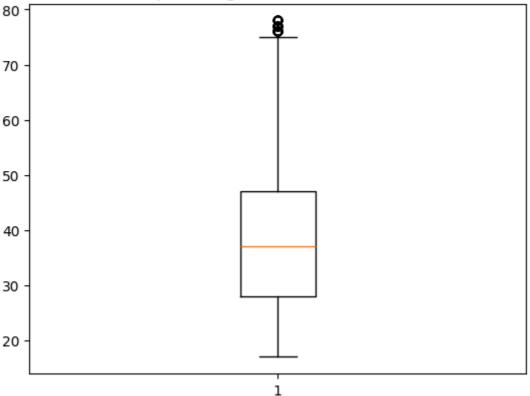
```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder, MinMaxScaler
from sklearn.impute import SimpleImputer
from sklearn.neighbors import KNeighborsClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.svm import SVC
from sklearn.neural_network import MLPClassifier
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_scc
# Load dataset
data = pd.read_csv("/content/adult 3.csv")
# Replace unknown values and remove unwanted rows
data['workclass'].replace({'?': 'others'}, inplace=True)
data['occupation'].replace({'?': 'others'}, inplace=True)
data = data[data['workclass'] != 'Without-pay']
data = data[data['workclass'] != 'Never-worked']
data = data[~data['education'].isin(['5th-6th', '1st-4th', 'Preschool'])]
# Drop redundancy
if 'education' in data.columns:
    data.drop(columns='education', inplace=True)
# Remove Outliers in 'age'
Q1 = data['age'].quantile(0.25)
Q3 = data['age'].quantile(0.75)
IQR = Q3 - Q1
lower bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR
data = data[(data['age'] >= lower_bound) & (data['age'] <= upper_bound)]</pre>
# Boxplot to visualize
plt.boxplot(data['age'])
plt.title("Boxplot of Age after Outlier Removal")
plt.show()
# Encode Categorical Columns
encoder = LabelEncoder()
for col in ['workclass', 'marital-status', 'occupation', 'relationship', 'race',
    data[col] = encoder.fit_transform(data[col])
# Encode Target
if 'income' in data.columns:
    data['income'] = encoder.fit transform(data['income']) # 0: <=50K, 1: >50K
# Split features and labels
X = data.drop(columns='income')
y = data['income']
```

```
# Impute missing values
imputer = SimpleImputer(strategy='mean')
X_imputed = imputer.fit_transform(X)
# MinMax Scaling
scaler = MinMaxScaler()
X_scaled = scaler.fit_transform(X_imputed)
# Train-Test Split
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2, r
# Apply Models and Evaluate
models = {
    "KNN": KNeighborsClassifier(),
    "Logistic Regression": LogisticRegression(max_iter=1000),
    "Decision Tree": DecisionTreeClassifier(),
    "SVC": SVC(probability=True),
    "MLP (ANN)": MLPClassifier(max_iter=1000, random_state=42)
}
```

Boxplot of Age after Outlier Removal



```
results = []
for name, model in models.items():
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)
    metrics = {
        "Model": name,
        "Accuracy": accuracy_score(y_test, y_pred),
        "Precision": precision_score(y_test, y_pred),
```

```
"Recall": recall_score(y_test, y_pred),
        "F1 Score": f1 score(y test, y pred)
    results.append(metrics)
    print(f"\n=== {name} Classification Report ===")
    print(classification_report(y_test, y_pred))
# Visualize Results
results_df = pd.DataFrame(results)
# Bar Graph Comparison
results_df.set_index("Model")[['Accuracy', 'Precision', 'Recall', 'F1 Score']].plot(kind=
plt.title("Model Comparison - Performance Metrics")
plt.ylabel("Score")
plt.ylim(0, 1)
plt.grid(axis='y')
plt.tight_layout()
plt.show()
# Print Final Comparison Table
print("\n=== Final Comparison Table ===")
print(results_df.round(4))
```



1011 61					
=== KNN Clas			C4		
	precision	recall	f1-score	support	
0	0.86	0.91	0.89	7185	
1	0.67	0.57	0.62	2369	
1	0.67	0.57	0.62	2309	
accuracy			0.82	9554	
macro avg	0.77	0.74	0.75	9554	
weighted avg	0.82	0.82	0.82	9554	
weighted avg	0.02	0.02	0.02	JJJ-	
=== Logistic	Regression	Classifica	ation Repor	t ===	
· ·	precision		f1-score	support	
	•			• • •	
0	0.84	0.94	0.89	7185	
1	0.72	0.46	0.56	2369	
accuracy			0.82	9554	
macro avg	0.78	0.70	0.72	9554	
weighted avg	0.81	0.82	0.81	9554	
=== Decision	Tree Class:	ification F	Report ===		
	precision	recall	f1-score	support	
0	0.88	0.87	0.87	7185	
1	0.62	0.62	0.62	2369	
accuracy			0.81	9554	
macro avg	0.75	0.75	0.75	9554	
weighted avg	0.81	0.81	0.81	9554	
C) (C, C]					
=== SVC Clas		•	C4		
	precision	recall	f1-score	support	
0	0.00	0.05	0.00	71.05	
0	0.86	0.95	0.90	7185	
1	0.77	0.53	0.63	2369	
accunacy			0.84	9554	
accuracy	0.82	0.74	0.84	9554 9554	
macro avg		0.74	0.76		
weighted avg	0.84	v.04	۵.03	9554	
=== MLP (ANN) Classific:	ation Renor	rt ===		
TIEL (MIN	precision	recall	f1-score	support	
	p. cc131011	, ccarr	. 2 30010	Suppor C	
0	0.87	0.94	0.90	7185	
1	0.77	0.57	0.66	2369	
_	J.77	0.57	0.00	2505	
accuracy			0.85	9554	
macro avg	0.82	0.76	0.78	9554	
weighted avg	0.84	0.85	0.84	9554	
- 0		3.03			
		Mod	del Comparison - P	erformance Metrics	5
1.0					

import pandas as pd

Accuracy

```
from sklearn.svm import SVC
from sklearn.model selection import train test split
import joblib
# Load data
data = pd.read_csv("/content/adult 3.csv")
# Basic cleaning
data['workclass'].replace({'?': 'others'}, inplace=True)
data['occupation'].replace({'?': 'others'}, inplace=True)
data = data[data['workclass'] != 'Without-pay']
data = data[data['workclass'] != 'Never-worked']
data = data[~data['education'].isin(['5th-6th', '1st-4th', 'Preschool'])]
if 'education' in data.columns:
    data.drop(columns='education', inplace=True)
# Encode categorical variables
le_occupation = LabelEncoder()
data['occupation'] = le_occupation.fit_transform(data['occupation'].astype(str))
# Target encoding
target_le = LabelEncoder()
data['income'] = target_le.fit_transform(data['income'])
# Select features
features = ['age', 'hours-per-week', 'occupation'] # 'education' removed
X = data[features]
y = data['income']
# Scale
scaler = MinMaxScaler()
X_scaled = scaler.fit_transform(X)
# Train model
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2, random_st
model = SVC(probability=True)
model.fit(X train, y train)
# Save model and encoders
joblib.dump(model, "salary model.pkl")
joblib.dump(le_occupation, "le_occupation.pkl")
joblib.dump(scaler, "scaler.pkl")
\rightarrow \uparrow /tmp/ipython-input-11-3688869629.py:11: FutureWarning: A value is trying to be set on
     The behavior will change in pandas 3.0. This inplace method will never work because t
     For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({
       data['workclass'].replace({'?': 'others'}, inplace=True)
     /tmp/ipython-input-11-3688869629.py:12: FutureWarning: A value is trying to be set on
     The behavior will change in pandas 3.0. This inplace method will never work because t
     For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({
```

```
data['occupation'].replace({'?': 'others'}, inplace=True)
     ['scaler.pkl']
loaded model = joblib.load("salary model.pkl")
loaded_occupation_encoder = joblib.load("le_occupation.pkl")
loaded scaler = joblib.load("scaler.pkl")
def predict_income(age, hours_per_week, education, occupation):
    # Create a DataFrame from input
    input data = pd.DataFrame({
        'age': [age],
        'hours-per-week': [hours_per_week],
        'education': [education],
        'occupation': [occupation]
    })
    # Encode occupation - Handle potential unseen values
        input_data['occupation'] = loaded_occupation_encoder.transform(input_data['occupa
    except ValueError:
        # Handle unseen occupation by assigning a default value or raising an error
        # For simplicity, let's assign a value outside the range of encoded values
        input_data['occupation'] = -1 # Or handle as needed
    # Select and order features for scaling
    features_for_scaling = input_data[['age', 'hours-per-week', 'occupation']]
    # Scale numerical features
    scaled_input = loaded_scaler.transform(features_for_scaling)
    # Predict using the loaded model
    prediction = loaded_model.predict(scaled_input)
    # Convert numerical prediction back to human-readable label
    # Assuming the target encoder was fitted on ['<=50K', '>50K'] resulting in 0 and 1
    income label = "<=50K" if prediction[0] == 0 else ">50K"
    return income_label
import gradio as gr
# Define input components
age_input = gr.Number(label="Age", minimum=0, maximum=100)
hours_input = gr.Number(label="Hours per Week", minimum=0, maximum=168)
# Since the model only used occupation, education input will not be used in prediction
# but we will include it in the interface as per the original request.
education_input = gr.Textbox(label="Education Level")
occupation_input = gr.Textbox(label="Occupation")
```

```
# Define output component
output_text = gr.Textbox(label="Predicted Income")

# Create the Gradio interface
iface = gr.Interface(
    fn=predict_income,
    inputs=[age_input, hours_input, education_input, occupation_input],
    outputs=output_text,
    title="Income Prediction Model",
    description="Predict income based on age, hours worked per week, education, and occup
)

iface.launch(share=True)

Rerunning server... use `close()` to stop if you need to change `launch()` parameters
    Colab notebook detected. To show errors in colab notebook, set debug=True in launch()
```

This share link expires in 1 week. For free permanent hosting and GPU upgrades, run `

Income Prediction Model

Predict income based on age, hours worked per week, education, and occupation.

* Running on public URL: https://eab4500e2351dfafc1.gradio.live

Age	
0	
Hours per Week	
0	
Education Level	
Occupation	

Clear Submit

```
import gradio as gr

# Step 1: Login UI
def login(username, password):
    if username == "deeksha" and password == "yaminn":
```

```
return gr.update(visible=True), "Login Successful!"
    else:
        return gr.update(visible=False), "Login Failed. Try again."
# Step 2: Salary Prediction App (hidden until login success)
with gr.Blocks() as demo:
    gr.Markdown("## i Login Now")
    username = gr.Textbox(label="Username")
    password = gr.Textbox(label="Password", type="password")
    login_btn = gr.Button("LOGIN")
    login_msg = gr.Textbox(label="", interactive=False)
    # Salary Prediction Section (Initially Hidden)
    with gr.Column(visible=False) as predict_ui:
        gr.Markdown("Salary Prediction")
        age = gr.Number(label="Age")
        hours = gr.Number(label="Hours per Week")
        edu = gr.Textbox(label="Education") # Note: This input is not used in the predic
        occ = gr.Textbox(label="Occupation")
        predict_btn = gr.Button("Predict")
        output = gr.Textbox(label="Predicted Income")
        # Link the predict button to the predict_income function with the correct inputs
        predict_btn.click(fn=predict_income, inputs=[age, hours, edu, occ], outputs=outputs=
    login_btn.click(fn=login, inputs=[username, password], outputs=[predict_ui, login_ms
    It looks like you are running Gradio on a hosted a Jupyter notebook. For the Gradic a
     Colab notebook detected. To show errors in colab notebook, set debug=True in launch )
     * Running on public URL: <a href="https://5454e50959ce546d2f.gradio.live">https://5454e50959ce546d2f.gradio.live</a>
                                                                               grades, run `
           Login Now
            Username
            Password
```

LOGIN